

STRUCTURED DATA

Accessing structured data with SQL is quite different from the full text search of documents on the Web. Structured data in the relational model means data that can be represented in tables -- rows and columns. Each row in a table represents a different object, and the columns represent various "attributes" of the object. The columns have names and integrity constraints that specify valid values.

Since the column values are named and are represented in a consistent format, you can select rows very precisely, based on their contents. This is especially helpful in dealing with numeric data. You can also join together data from different tables, based on matching column values. You can do useful types of analysis, listing objects that are in one table and missing (or present, or have specific attributes) from a related table. You can extract from a large table precisely those rows of interest, regrouping them and generating simple statistics on them.

This document contains examples of:

- [creating a table](#)
- [creating a view](#)
- [inserting rows](#)
- [updating rows](#)
- [deleting rows](#)
- [commit -- making changes permanent](#)
- [rollback -- undoing temporary changes](#)
- [enforcing integrity constraints](#)
- [using an Embedded C program](#)

INTERACTIVE SQL EXAMPLES

create a table to store information about weather observation stations:

-- No duplicate ID fields allowed

```
CREATE TABLE STATION
(ID INTEGER PRIMARY KEY,
CITY CHAR(20),
STATE CHAR(2),
LAT_N REAL,
LONG_W REAL);
```

populate the table STATION with a few rows:

```
INSERT INTO STATION VALUES (13, 'Phoenix', 'AZ', 33, 112);
INSERT INTO STATION VALUES (44, 'Denver', 'CO', 40, 105);
INSERT INTO STATION VALUES (66, 'Caribou', 'ME', 47, 68);
```

query to look at table STATION in undefined order:

```
SELECT * FROM STATION;
```

ID	CITY	STATE	LAT_N	LONG_W
13	Phoenix	AZ	33	112
44	Denver	CO	40	105
66	Caribou	ME	47	68

query to select Northern stations (Northern latitude > 39.7):

-- selecting only certain rows is called a "restriction".

```
SELECT * FROM STATION
```

```
WHERE LAT_N > 39.7;
```

ID	CITY	STATE	LAT_N	LONG_W
44	Denver	CO	40	105
66	Caribou	ME	47	68

query to select only ID, CITY, and STATE columns:

-- selecting only certain columns is called a "projection".

```
SELECT ID, CITY, STATE FROM STATION;
```

ID	CITY	STATE
13	Phoenix	AZ
44	Denver	CO
66	Caribou	ME

query to both "restrict" and "project":

```
SELECT ID, CITY, STATE FROM STATION
```

```
WHERE LAT_N > 39.7;
```

ID	CITY	STATE
44	Denver	CO
66	Caribou	ME

create another table to store normalized temperature and precipitation data:

-- ID field must match some STATION table ID

(so name and location will be known).

-- allowable ranges will be enforced for other values.

-- no duplicate ID and MONTH combinations.

-- temperature is in degrees Fahrenheit.

-- rainfall is in inches.

```
CREATE TABLE STATS
```

```
(ID INTEGER REFERENCES STATION(ID),
```

```
MONTH INTEGER CHECK (MONTH BETWEEN 1 AND 12),
```

```
TEMP_F REAL CHECK (TEMP_F BETWEEN -80 AND 150),
```

```
RAIN_I REAL CHECK (RAIN_I BETWEEN 0 AND 100),
```

```
PRIMARY KEY (ID, MONTH));
```

populate the table STATS with some statistics for January and July:

```

INSERT INTO STATS VALUES (13, 1, 57.4, 0.31);
INSERT INTO STATS VALUES (13, 7, 91.7, 5.15);
INSERT INTO STATS VALUES (44, 1, 27.3, 0.18);
INSERT INTO STATS VALUES (44, 7, 74.8, 2.11);
INSERT INTO STATS VALUES (66, 1, 6.7, 2.10);
INSERT INTO STATS VALUES (66, 7, 65.8, 4.52);

```

query to look at table STATS in undefined order:

```
SELECT * FROM STATS;
```

ID	MONTH	TEMP_F	RAIN_I
13	1	57.4	.31
13	7	91.7	5.15
44	1	27.3	.18
44	7	74.8	2.11
66	1	6.7	2.1
66	7	65.8	4.52

query to look at table STATS, picking up location information by joining with table STATION on the ID column:

```

-- matching two tables on a common column is called a "join".
-- the column names often match, but this is not required.
-- only the column values are required to match.

```

```

SELECT * FROM STATION, STATS
WHERE STATION.ID = STATS.ID;

```

ID	CITY	ST	LAT_N	LONG_W	ID	MONTH	TEMP_F	RAIN_I
13	Phoenix	AZ	33	112	13	1	57.4	.31
13	Phoenix	AZ	33	112	13	7	91.7	5.15
44	Denver	CO	40	105	44	1	27.3	.18
44	Denver	CO	40	105	44	7	74.8	2.11
66	Caribou	ME	47	68	66	1	6.7	2.1
66	Caribou	ME	47	68	66	7	65.8	4.52

query to look at the table STATS, ordered by month and greatest rainfall, with columns rearranged:

```

SELECT MONTH, ID, RAIN_I, TEMP_F
FROM STATS
ORDER BY MONTH, RAIN_I DESC;

```

MONTH	ID	RAIN_I	TEMP_F
1	66	2.1	6.7
1	13	.31	57.4
1	44	.18	27.3

7	13	5.15	91.7
7	66	4.52	65.8
7	44	2.11	74.8

query to look at temperatures for July from table STATS, lowest temperatures first, picking up city name and latitude by joining with table STATION on the ID column:

```
SELECT LAT_N, CITY, TEMP_F
FROM STATS, STATION
WHERE MONTH = 7
AND STATS.ID = STATION.ID
ORDER BY TEMP_F;
```

LAT_N	CITY	TEMP_F
47	Caribou	65.8
40	Denver	74.8
33	Phoenix	91.7

query to show MAX and MIN temperatures as well as average rainfall for each station:

```
SELECT MAX(TEMP_F), MIN(TEMP_F), AVG(RAIN_I), ID
FROM STATS
GROUP BY ID;
```

MAX(TEMP_F)	MIN(TEMP_F)	AVG(RAIN_I)	ID
91.7	57.4	2.73	13
74.8	27.3	1.145	44
65.8	6.7	3.31	66

query (with subquery) to show stations with year-round average temperature above 50 degrees:

-- rows are selected from the STATION table based on related values in the STATS table.

```
SELECT * FROM STATION
WHERE 50 < (SELECT AVG(TEMP_F) FROM STATS
WHERE STATION.ID = STATS.ID);
```

ID	CITY	ST	LAT_N	LONG_W
13	Phoenix	AZ	33	112
44	Denver	CO	40	105

create a view (derived table or persistent query) to convert Fahrenheit to Celsius and inches to centimeters:

```
CREATE VIEW METRIC_STATS (ID, MONTH, TEMP_C, RAIN_C) AS
SELECT ID,
MONTH,
(TEMP_F - 32) * 5 / 9,
RAIN_I * 0.3937
FROM STATS;
```

query to look at table STATS in a metric light (through the new view):

```
SELECT * FROM METRIC_STATS;
```

ID	MONTH	TEMP_C	RAIN_C
13	1	14.111111	.122047
13	7	33.166667	2.027555
44	1	-2.611111	.070866
44	7	23.777778	.830707
66	1	-14.055556	.82677
66	7	18.777778	1.779524

another metric query restricted to January below-freezing (0 Celsius) data, sorted on rainfall:

```
SELECT * FROM METRIC_STATS
WHERE TEMP_C < 0 AND MONTH = 1
ORDER BY RAIN_C;
```

ID	MONTH	TEMP_C	RAIN_C
44	1	-2.611111	.070866
66	1	-14.055556	.82677

Interactive SQL Update Examples

update all rows of table STATS to compensate for faulty rain gauges known to read 0.01 inches low:

```
UPDATE STATS SET RAIN_I = RAIN_I + 0.01;
```

and take a look:

```
SELECT * FROM STATS;
```

ID	MONTH	TEMP_F	RAIN_I
13	1	57.4	.32
13	7	91.7	5.16
44	1	27.3	.19
44	7	74.8	2.12

66	1	6.7	2.11
66	7	65.8	4.53

update one row, Denver's July temperature reading, to correct a data entry error:

```
UPDATE STATS SET TEMP_F = 74.9
WHERE ID = 44
AND MONTH = 7;
```

and take a look:

```
SELECT * FROM STATS;
```

ID	MONTH	TEMP_F	RAIN_I
13	1	57.4	.32
13	7	91.7	5.16
44	1	27.3	.19
44	7	74.9	2.12
66	1	6.7	2.11
66	7	65.8	4.53

make the above changes permanent:

-- they were only temporary until now.

```
COMMIT WORK;
```

update two rows, Denver's rainfall readings:

```
UPDATE STATS SET RAIN_I = 4.50
WHERE ID = 44;
```

and take a look:

```
SELECT * FROM STATS;
```

ID	MONTH	TEMP_F	RAIN_I
13	1	57.4	.32
13	7	91.7	5.16
44	1	27.3	4.5
44	7	74.9	4.5
66	1	6.7	2.11
66	7	65.8	4.53

Oops! We meant to update just the July reading! Undo that update:

-- undoes only updates since the last COMMIT WORK.

ROLLBACK WORK;

and take a look:

SELECT * FROM STATS;

ID	MONTH	TEMP_F	RAIN_I
13	1	57.4	.32
13	7	91.7	5.16
44	1	27.3	.19
44	7	74.9	2.12
66	1	6.7	2.11
66	7	65.8	4.53

now update Denver's July rainfall reading and make it permanent:

UPDATE STATS SET RAIN_I = 4.50

WHERE ID = 44

AND MONTH = 7;

COMMIT WORK;

and take a look:

SELECT * FROM STATS;

ID	MONTH	TEMP_F	RAIN_I
13	1	57.4	.32
13	7	91.7	5.16
44	1	27.3	.19
44	7	74.9	4.5
66	1	6.7	2.11
66	7	65.8	4.53

delete July data and East Coast data from both tables:

-- note that we use longitude values from the related STATION table to determine which STAT stations were east of 90 degrees.

DELETE FROM STATS

WHERE MONTH = 7

OR ID IN (SELECT ID FROM STATION

WHERE LONG_W < 90);

```
DELETE FROM STATION WHERE LONG_W < 90;
```

```
COMMIT WORK;
```

and take a look:

```
SELECT * FROM STATION;
```

ID	CITY	ST	LAT_N	LONG_W
13	Phoenix	AZ	33	112
44	Denver	CO	40	105

```
SELECT * FROM STATS;
```

ID	MONTH	TEMP_F	RAIN_I
13	1	57.4	.32
44	1	27.3	.19

View METRIC_STATS, a Fahrenheit-to-Centigrade and inches-to-centimeters conversion of table STATS, reflects the updates made to the underlying table.

```
SELECT * FROM METRIC_STATS;
```

ID	MONTH	TEMP_C	RAIN_C
13	1	14.111111	.125984
44	1	-2.611111	.074803

SQL Constraints

SQL enforces data integrity constraints.

Attempt to insert a row for an unknown observation station.

```
-- The ID value of 33 does not match a station ID value in the STATION table.
```

```
-- This is a violation of referential integrity.
```

```
INSERT INTO STATS VALUES (33,8,27.4,.19);
```

error message
violation of constraint STATS_FOREIGN1 caused operation to fail

Attempt to update a row with a temperature below the range -80 TO 150.

```
UPDATE STATS SET TEMP_F = -100 WHERE ID = 44 AND MONTH = 1;
```

error message
violation of constraint STATS_CHECK2 caused operation to fail

Attempt to insert a row with negative rainfall measurement, outside the range 0 to 100.

```
INSERT INTO STATS VALUES (44,8,27.4,-.03);
```

error message

violation of constraint STATS_CHECK3 caused operation to fail

Attempt to insert a row with month 13, outside the range of 1 to 12.

```
INSERT INTO STATS VALUES (44,13,27.4,.19);
```

error message

violation of constraint STATS_CHECK1 caused operation to fail

Attempt to insert a row with a temperature above the range -80 TO 150.

```
INSERT INTO STATS VALUES (44,8,160,.19);
```

error message

violation of constraint STATS_CHECK2 caused operation to fail

Attempt to insert a row with no constraint violations.

```
INSERT INTO STATS VALUES (44,8,27.4,.10);
```

status message

1 row inserted

and take a look:

```
SELECT * FROM STATS;
```

ID	MONTH	TEMP_F	RAIN_I
44	8	27.4	.10
13	1	57.4	.32
44	1	27.3	.19

Attempt to insert a second row of August statistics for station 44.

-- This is a violation of the primary key constraint.

-- Only one row for each station and month combination is allowed.

```
INSERT INTO STATS VALUES (44,8,160,.19);
```

error message

violation of constraint STATS_PRIMARY_ID_MONTH caused operation to fail

Embedded SQL C Program Example

Embedded C program to do the following:

Starting with a station name (Denver, in this example), look up the station ID.

Print the station ID.

List all rows for that station ID.

- shows single-row select and use of cursor
- note that all C-language variables used in SQL statements are declared in the DECLARE SECTION.
- note that all SQL statements begin with the syntax EXEC SQL and terminate with a semicolon.

```
#include<stdio.h>
#include<string.h>
EXEC SQL BEGIN DECLARE SECTION;

    long station_id;
    long mon;
    float temp;
    float rain;
    char city_name[21];
    long SQLCODE;
EXEC SQL END DECLARE SECTION;
main()
{
    /* the CONNECT statement, if needed, goes here */
    strcpy(city_name,"Denver");
    EXEC SQL SELECT ID INTO :station_id
        FROM STATION
        WHERE CITY = :city_name;
    if (SQLCODE == 100)
    {
        printf("There is no station for city %s\n",city_name);
        exit(0);
    }
    printf("For the city %s, Station ID is %ld\n",city_name,station_id);
    printf("And here is the weather data:\n");
    EXEC SQL DECLARE XYZ CURSOR FOR
        SELECT MONTH, TEMP_F, RAIN_I
        FROM STATS
        WHERE ID = :station_id
        ORDER BY MONTH;
    EXEC SQL OPEN XYZ;
    while (SQLCODE != 100) {
        EXEC SQL FETCH XYZ INTO :mon, :temp, :rain;
        if (SQLCODE == 100)
```

```
        printf("end of list\n");
    else
        printf("month = %ld, temperature = %f, rainfall = %f\n",mon,temp,rain);
    }
EXEC SQL CLOSE XYZ;
exit(0);
}
```

Execution log:

For the city Denver, Station ID is 44
And here is the weather data:
month = 1, temperature = 27.299999, rainfall = 0.180000
month = 7, temperature = 74.800003, rainfall = 2.110000
end of list

[return to SQL Table of Contents](#)